



PyTorch vs Triton Inference Server

OpenCLIP ViT-B-32 on RTX 3070


Single-Image Latency



Batch Processing Throughput



Concurrent Throughput (16 workers, 200 requests)



Cost to Process 10,000 Images (@ \$0.20/hour GPU)

